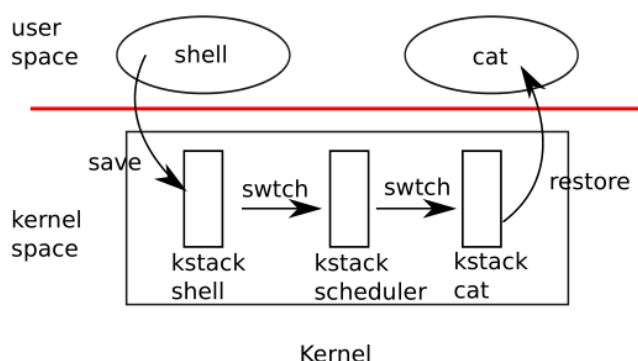


## Lab7: Multithreading 实验报告

线程调度过程如下图所示：



线程调度的过程主要是保存 **context** 上下文状态，因为这里的切换全都是以函数调用的形式，因此这里只需要保存被调用者保存的寄存器（Callee-saved register）即可，调用者的寄存器会自动保存。

1. 首先切换分支：

```
1 $ git fetch
2 $ git checkout thread
3 $ make clean
```

## Uthread: switching between threads (moderate)

实现一个用户线程调度的方法。

这里的“线程”是完全用户态实现的，多个线程也只能运行在一个 CPU 上，并且没有时钟中断来强制执行调度，需要线程函数本身在合适的时候主动 yield 释放 CPU。

1. 上下文切换：

借鉴 `kernel/swtch.S`，在 `user/uthread_switch.S` 中需要实现上下文切换的代码

调用者寄存器被c编译器保存在thread\_schedule的堆栈中

```

9  ~ thread_switch:
10     /* YOUR CODE HERE */
11     sd ra, 0(a0)
12     sd sp, 8(a0)
13     sd s0, 16(a0)
14     sd s1, 24(a0)
15     sd s2, 32(a0)
16     sd s3, 40(a0)
17     sd s4, 48(a0)
18     sd s5, 56(a0)
19     sd s6, 64(a0)
20     sd s7, 72(a0)
21     sd s8, 80(a0)
22     sd s9, 88(a0)
23     sd s10, 96(a0)
24     sd s11, 104(a0)
25
26     ld ra, 0(a1)
27     ld sp, 8(a1)
28     ld s0, 16(a1)
29     ld s1, 24(a1)
30     ld s2, 32(a1)
31     ld s3, 40(a1)
32     ld s4, 48(a1)
33     ld s5, 56(a1)
34     ld s6, 64(a1)
35     ld s7, 72(a1)
36     ld s8, 80(a1)
37     ld s9, 88(a1)
38     ld s10, 96(a1)
39     ld s11, 104(a1)
40
41     ret    /* return to ra */

```

## 2. 定义上下文字段:

借鉴 *kernel/proc.h* 中的上下文结构体, 在 *user/uthread.c* 中定义一个 `context` 结构体, 用于保存 ra、sp 以及 callee-saved registers, 并加入到 `thread` 结构体中。

```

13 // Saved registers for user context switches.
14 struct context {
15     uint64 ra;
16     uint64 sp;
17
18     // callee-saved
19     uint64 s0;
20     uint64 s1;
21     uint64 s2;
22     uint64 s3;
23     uint64 s4;
24     uint64 s5;
25     uint64 s6;
26     uint64 s7;
27     uint64 s8;
28     uint64 s9;
29     uint64 s10;
30     uint64 s11;
31 };
32
33 struct thread {
34     char    stack[STACK_SIZE]; /* the thread's stack */
35     int     state;              /* FREE, RUNNING, RUNNABLE */
36     struct context context;     /* 存储线程上下文切换时保存的寄存器 */
37 };

```

## 3. 在 *kernel/uthread.c* 中的 `thread_schedule` 函数中调用 `thread_switch` 进行上下文切换

```

77     if (current_thread != next_thread) {           /* switch threads? */
78         next_thread->state = RUNNING;
79         t = current_thread;
80         current_thread = next_thread;
81         /* YOUR CODE HERE
82          * Invoke thread_switch to switch from t to next_thread:
83          * thread_switch(??, ??);
84          */
85         thread_switch((uint64) &t->context, (uint64) &current_thread->context);
86     } else
87         next_thread = 0;
88 }

```

4. 补充 `kernel/uthread.c` 中的 `thread_create` 函数：将线程函数的入口地址保存到返回地址 `ra` 中，使得 `thread_switch` 最后返回到该地址，从而运行线程代码；将 `sp` 指向栈底（由于栈从高到低生长，所以 `sp` 应该指向栈的最高地址）

```

90 void
91 thread_create(void (*func)())
92 {
93     struct thread *t;
94
95     for (t = all_thread; t < all_thread + MAX_THREAD; t++) {
96         if (t->state == FREE) break;
97     }
98     t->state = RUNNABLE;
99     // YOUR CODE HERE
100    t->context.ra = (uint64)func;
101    t->context.sp = (uint64)&t->stack + (STACK_SIZE - 1);
102 }

```

5. 执行测试命令：`uthread`

```

220110512@comp6:~/xv6-labs-2020$ make qemu
qemu-system-riscv64 -machine virt -bios none -kernel kernel/kernel -m 128M -smp 3 -nographic -drive file=fs.img,if=none,format=raw,id=x0 -device virtio-blk-device,drive=x0,bus=virtio-mmio-bus.0

xv6 kernel is booting

hart 1 starting
hart 2 starting
init: starting sh
$ uthread
thread_a started
thread_b started
thread_c started
thread_c 0
thread_a 0
thread_b 0
thread_c 1
thread_a 1
thread_b 1
thread_c 2
thread_a 2
thread_b 2
thread_c 3
thread_a 3
thread_b 3
thread_c 4

```

```

thread_a 96
thread_b 96
thread_c 97
thread_a 97
thread_b 97
thread_c 98
thread_a 98
thread_b 98
thread_c 99
thread_a 99
thread_b 99
thread_c: exit after 100
thread_a: exit after 100
thread_b: exit after 100
thread_schedule: no runnable threads
$

```

## Using threads (moderate)

1. 文件 `notxv6/ph.c` 包含一个简单的哈希表，如果单个线程使用，该哈希表是正确的，但是多个线程使用时，该哈希表是不正确的。

键入 `make ph`，`./ph 1` 和 `./ph 2` 可以看到，当有两个线程同时向哈希表中添加条目时，实现了较优的并行性能，但两个线程均丢失了16550个键，表明大量本应在哈希表中的键不存在在哈希表中

```

220110512@comp6:~/xv6-labs-2020$ make ph
gcc -o ph -g -O2 notxv6/ph.c -pthread
220110512@comp6:~/xv6-labs-2020$ ./ph 1
100000 puts, 11.695 seconds, 8551 puts/second
0: 0 keys missing
100000 gets, 12.050 seconds, 8299 gets/second
220110512@comp6:~/xv6-labs-2020$ htop
220110512@comp6:~/xv6-labs-2020$ ./ph 2
100000 puts, 5.637 seconds, 17739 puts/second
1: 16550 keys missing
0: 16550 keys missing
200000 gets, 12.814 seconds, 15608 gets/second

```

2. 回答问题：为什么多线程会导致键的缺失，而单键值不会？

设计这样一个序列：

[假设键 k1、k2 属于同个 bucket]

- thread 1: 尝试设置 k1
- thread 1: 发现 k1 不存在，尝试在 bucket 末尾插入 k1

---scheduler 切换到 thread 2---

- thread 2: 尝试设置 k2
- thread 2: 发现 k2 不存在，尝试在 bucket 末尾插入 k2
- thread 2: 分配 entry，在桶末尾插入 k2

---scheduler 切换回 thread 1---

- thread 1: 分配 entry, 没有意识到 k2 的存在, 在其认为的 “桶末尾” (实际为 k2 所处位置) 插入 k1

[k1 被插入, 但是由于被 k1 覆盖, k2 从桶中消失了, 引发了键值丢失]

由此造成了键值缺失

### 3. 利用加锁操作, 解决哈希表 race-condition 导致的数据丢失问题

如果只加一个锁, 锁的粒度很大, 会导致丢失性能, 结果还不如不加锁的单线程。因此需要将锁的粒度减小, 为每个槽位 (bucket) 加一个锁。

在 *notxv6/ph.c* 中定义 lock 锁, 并修改 put、get、main 函数, 在适当位置加锁和解锁:

```
20 pthread_mutex_t lock[NBUCKET];
```

```
40 static
41 void put(int key, int value)
42 {
43     int i = key % NBUCKET;
44
45     pthread_mutex_lock(&lock[i]);
46     // is the key already present?
47     struct entry *e = 0;
48     for (e = table[i]; e != 0; e = e->next) {
49         if (e->key == key)
50             break;
51     }
52     if(e){
53         // update the existing key.
54         e->value = value;
55     } else {
56         // the new is new.
57         insert(key, value, &table[i], table[i]);
58     }
59     pthread_mutex_unlock(&lock[i]);
60 }
```

```
62 static struct entry*
63 get(int key)
64 {
65     int i = key % NBUCKET;
66
67     pthread_mutex_lock(&lock[i]);
68     struct entry *e = 0;
69     for (e = table[i]; e != 0; e = e->next) {
70         if (e->key == key) break;
71     }
72     pthread_mutex_unlock(&lock[i]);
73     return e;
74 }
```

```

103 int
104 main(int argc, char *argv[])
105 {
106     pthread_t *tha;
107     void *value;
108     double t1, t0;
109
110     if (argc < 2) {
111         fprintf(stderr, "Usage: %s nthreads\n", argv[0]);
112         exit(-1);
113     }
114     nthread = atoi(argv[1]);
115     tha = malloc(sizeof(pthread_t) * nthread);
116     srand(0);
117     assert(NKEYS % nthread == 0);
118     for (int i = 0; i < NKEYS; i++) {
119         keys[i] = random();
120     }
121     for (int i = 0; i < NBUCKET; i++){
122         pthread_mutex_init(&lock[i], NULL);
123     }

```

4. 键入 `make ph` , `./ph 1` 和 `./ph 2` :

```

220110512@comp6:~/xv6-labs-2020$ make ph
gcc -o ph -g -O2 notxv6/ph.c -pthread
220110512@comp6:~/xv6-labs-2020$ ./ph 1
100000 puts, 11.910 seconds, 8396 puts/second
220110512@comp6:~/xv6-labs-2020$ ./ph 2
100000 puts, 6.822 seconds, 14658 puts/second
0: 0 keys missing
1: 0 keys missing
200000 gets, 13.384 seconds, 14943 gets/second
220110512@comp6:~/xv6-labs-2020$

```

实现了较优的并行性能，且不存在键值缺失

## Barrier(moderate)

实验实现思是：加锁，然后判断到达屏障点的线程数，如果所有线程都已到达就调用 `pthread_cond_broadcast` 接口唤醒其他线程，否则就调用 `pthread_cond_wait` 进行等待。

1. 在 `notxv6/barrier.c` 文件中完成 `barrier` 函数：

采用生产者消费者模式，如果还有线程没到达，就加入到队列中，等待唤起；如果最后一个线程到达了，就将轮数加一，然后唤醒所有等待这个条件变量的线程

```

25 static void
26 barrier()
27 {
28     // YOUR CODE HERE
29     //
30     // Block until all threads have called barrier() and
31     // then increment bstate.round.
32     //
33     pthread_mutex_lock(&bstate.barrier_mutex);
34     if(++bstate.nthread == nthread) {
35         // 记录等待进程数的计数器清零
36         bstate.nthread = 0;
37         // 屏障轮数加一
38         bstate.round++;
39         // 唤醒等待在此点上的所有进程
40         pthread_cond_broadcast(&bstate.barrier_cond);
41     } else {
42         // 等待
43         pthread_cond_wait(&bstate.barrier_cond, &bstate.barrier_mutex);
44     }
45     pthread_mutex_unlock(&bstate.barrier_mutex);
46 }

```

2. 键入 `make barrier` , `./barrier 1` 、 `./barrier 2` 和 `./barrier 4` :

```

220110512@comp6:~/xv6-labs-2020$ make barrier
gcc -o barrier -g -O2 notxv6/barrier.c -pthread
220110512@comp6:~/xv6-labs-2020$ ./barrier 1
OK; passed
220110512@comp6:~/xv6-labs-2020$ ./barrier 2
OK; passed
220110512@comp6:~/xv6-labs-2020$ ./barrier 4
OK; passed
220110512@comp6:~/xv6-labs-2020$

```

## 结果截图

执行命令 `make grade`

```

make[1]: Leaving directory '/home/students/220110512/xv6-labs-2020'
== Test uthread ==
$ make qemu-gdb
uthread: OK (4.7s)
== Test answers-thread.txt == answers-thread.txt: OK
== Test ph_safe == make[1]: Entering directory '/home/students/220110512/xv6-labs-2020'
make[1]: 'ph' is up to date.
make[1]: Leaving directory '/home/students/220110512/xv6-labs-2020'
ph_safe: OK (20.5s)
== Test ph_fast == make[1]: Entering directory '/home/students/220110512/xv6-labs-2020'
make[1]: 'ph' is up to date.
make[1]: Leaving directory '/home/students/220110512/xv6-labs-2020'
ph_fast: OK (43.2s)
== Test barrier == make[1]: Entering directory '/home/students/220110512/xv6-labs-2020'
make[1]: 'barrier' is up to date.
make[1]: Leaving directory '/home/students/220110512/xv6-labs-2020'
barrier: OK (3.9s)
== Test time ==
time: OK
Score: 60/60
220110512@comp6:~/xv6-labs-2020$

```